


Data Analysis

Practice 5: Intro to TF2&Keras

Dr. Nataliya K. Sakhnenko

Deep Learning Frameworks



Keras API

- TensorFlow is an industry-strength numerical computing framework that can run on CPU, GPU, or TPU. It can automatically compute the gradient of any differentiable expression, it can be distributed to many devices, and it can export programs to various external runtimes.
- Keras is the standard API to do deep learning with TensorFlow.
- The central class of Keras is the **Layer**. A layer encapsulates some weights and some Layer computation. Layers are assembled into models.
- Before you start training a model, you need to pick an optimizer, a loss, and some metrics, which you specify via the **model.compile()** method.
- To train a model, you can use the method, which runs mini-batch gradient descent **fit()** for you. You can also use it to monitor your loss and metrics on "validation data", a set of inputs that the model doesn't see during training.
- Once your model is trained, use the **model.predict()** method to generate predictions on new inputs.

Keras API

The Sequential API

```
import keras
from keras import layers

model = keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

The **Sequential model**, the most approachable API—it's basically a Python list. As such, it's limited to simple stacks of layers

The functional API

```
import keras
from keras import layers

inputs = keras.Input(shape=(10,))
x = layers.Dense(20, activation='relu')(x)
x = layers.Dense(20, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = keras.Model(inputs, outputs)
```

<https://keras.io/>

Recap: softmax

$$T: X \rightarrow Y, Y = \{1, 2, \dots, K\}$$

$$h_{\theta}(x) = \begin{Bmatrix} P(y=1|x; \theta) \\ P(y=2|x; \theta) \\ \dots \\ P(y=K|x; \theta) \end{Bmatrix} = \frac{1}{\sum_{i=1}^K \exp(\theta_i^T x)} \begin{Bmatrix} \exp(\theta_1^T x) \\ \exp(\theta_2^T x) \\ \dots \\ \exp(\theta_K^T x) \end{Bmatrix}$$

One-hot encoding

$$\begin{bmatrix} 1 \\ 0 \\ \dots \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ \dots \\ 0 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 0 \\ \dots \\ 1 \end{bmatrix}, Y \in R^K$$

1, 2, ..., K


Model Training APIs

https://keras.io/api/models/model_training_apis/	
Compile method	<code>Model.compile(optimizer="rmsprop", loss=None, metrics=None, ...)</code>
Fit method	<code>Model.fit(x=None, y=None, batch_size=None, epochs=1, verbose="auto", callbacks=None, validation_split=0.0, validation_data=None, ...)</code>
Evaluate method	<code>Model.evaluate(x=None, y=None, batch_size=None, verbose=1, ...)</code>
Predict method	<code>Model.predict(x, ...)</code>

- **x**: Input data. It could be e.g. a numpy array

Last-layer activation and Loss function

Problem type	Last-activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass classification	softmax	categorical_crossentropy
Regression	None	mse



Installation & compatibility

Keras comes packaged with TensorFlow 2 as `tensorflow.keras`. To start using Keras, simply [install TensorFlow 2](https://www.tensorflow.org/install) (<https://www.tensorflow.org/install>)

Keras/TensorFlow are compatible with:

- Python 3.7–3.10
- Ubuntu 16.04 or later
- Windows 7 or later
- macOS 10.12.6 (Sierra) or later (no GPU support)

MNIST - Logistic Regression

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
input_dim = 28 * 28 #784
```

```
x_train = x_train.reshape(60000, input_dim)
```

```
x_test = x_test.reshape(10000, input_dim)
```

```
x_train = x_train.astype('float32')
```

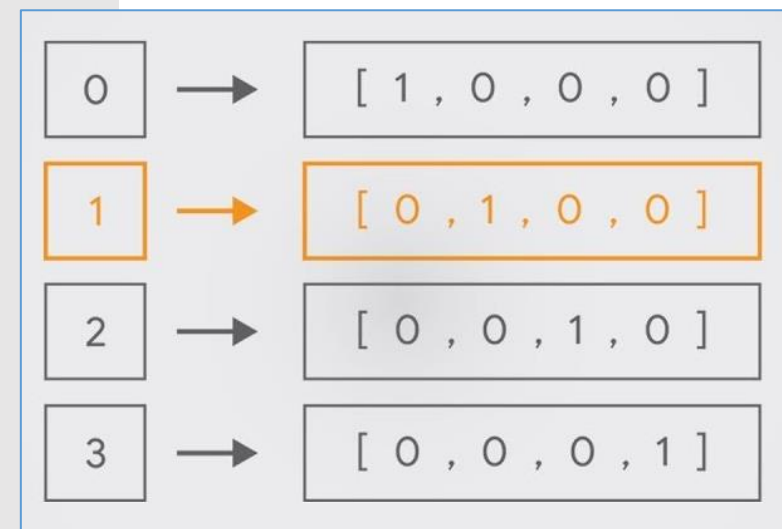

```
x_test = x_test.astype('float32')
```

```
x_train /= 255
```

```
x_test /= 255
```

```
y_train = keras.utils.to_categorical(y_train, num_classes)
```

```
y_test = keras.utils.to_categorical(y_test, num_classes)
```



One-hot encoding

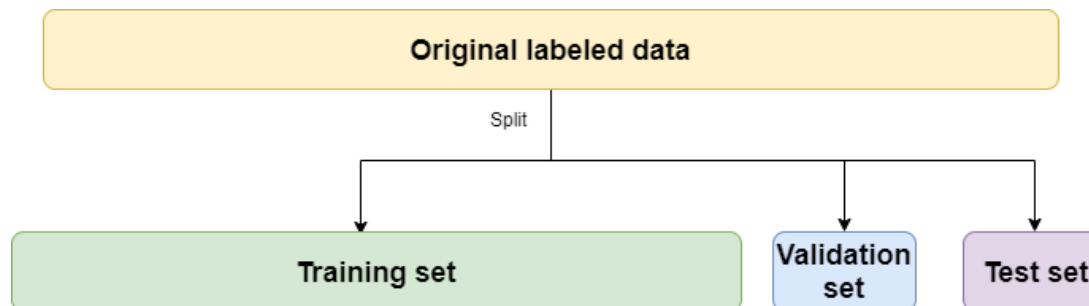
MNIST - Logistic Regression

```
model = keras.Sequential([layers.Dense(num_classes, activation="softmax")])  
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model.fit(x_train, y_train, epochs=10, batch_size=64, validation_split=0.2, verbose = 2)
```

Model summary

Layer (type)	Output shape	Param #
dense_10 (Dense)	(None, 10)	7850
Total params: 7,850		
Trainable params: 7,850		
Non-trainable params: 0		

```
score = model.evaluate(x_test, y_test)
```



MNIST - Multilayer FC ANN

```
model = keras.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 512)	401920
dense_13 (Dense)	(None, 512)	262656
dense_14 (Dense)	(None, 10)	5130

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

Callbacks my_callbacks = [tf.keras.callbacks.EarlyStopping(patience=2)]

Using Dropout

```
model = keras.Sequential()  
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))  
model.add(layers.Dropout(0.5))  
model.add(layers.Dense(512, activation='relu'))  
model.add(layers.Dropout(0.5))  
model.add(layers.Dense(10, activation='softmax'))  
  
model.summary()
```

Using Batch Normalization

```
model = keras.Sequential()

model.add(layers.Dense(512, input_shape=(784,)))
model.add(layers.BatchNormalization())
model.add(layers.Activation('relu'))
model.add(layers.Dropout(0.5))

model.add(layers.Dense(512))
model.add(layers.BatchNormalization())
model.add(layers.Activation('relu'))
model.add(layers.Dropout(0.5))

model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

MLP for text classification: Sentiment analysis, IMDb



<http://ai.stanford.edu/~amaas/data/sentiment/>



- ✓ We'll be using a dataset of 50,000 movie reviews taken from IMDb.
- ✓ The data is split evenly with 25k reviews intended for training and 25k for testing your classifier.
- ✓ Each set has 12.5k positive and 12.5k negative reviews.
- ✓ IMDb lets users rate movies on a scale from 1 to 10. To label these reviews the curator of the data labeled anything with ≤ 4 stars as negative and anything with ≥ 7 stars as positive. Reviews with 5 or 6 stars were left out.

SENTIMENT ANALYSIS



Discovering people opinions, emotions and feelings about
a product or service

Sentiment analysis (also known as **opinion mining** or **emotion AI**) refers to the use of NLP to systematically identify, extract, quantify, and study affective states and subjective information.

Sentiment Analysis, IMDB movies

```
from tensorflow.keras.datasets import imdb


(train_data, train_labels), (test_data, test_labels) = imdb.load_data()

x_train = vectorize_sequences(train_data)

model = keras.Sequential([
    layers.Dense(32, activation="relu"),
    layers.Dense(1, activation="sigmoid")])

model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

history = model.fit()
```





Bi-grams

Autoencoders


AE for dimensionality reduction

Train to reconstruct original data



An **autoencoder (AE)** is a type of ANN used to learn efficient data coding in an unsupervised manner. The goal of any AE is to reconstruct its own input. Usually, the AE first compresses the input into a smaller form, then transforms it back into an approximation of the input.

Denoising AE



- Reconstruction X' computed from the corrupted input X_{noise}
- Loss function compares X' reconstruction with the noiseless input X

AE: MNIST example

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
```

Load MNIST data

```
Input_dim = 784
(x_train, _), (x_test, _) = mnist.load_data()
x_train = x_train.reshape(60000, input_dim)
x_test = x_test.reshape(10000, input_dim)
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
```

```
def build_encoder():
    input_img = keras.Input(shape=(input_dim,))
    encoding_layer1 = layers.Dense(128, activation='relu')(input_img)
    encoding_layer2 = layers.Dense(64, activation='relu')(encoding_layer1)
    encoded = layers.Dense(latent_dim, activation='relu')(encoding_layer2)
    return keras.Model(input_img, encoded)

def build_decoder():
    decoder_input = keras.Input(shape=(latent_dim,))
    decoding_layer1 = layers.Dense(64, activation='relu')(decoder_input)
    decoding_layer2 = layers.Dense(128, activation='relu')(decoding_layer1)
    decoded = layers.Dense(input_dim, activation='sigmoid')(decoding_layer2)
    return keras.Model(decoder_input, decoded)
```


```
latent_dim = 32
encoder = build_encoder()
decoder = build_decoder()
```

```
input_images = keras.Input(shape=(input_dim,))
encoded_repr = encoder(input_images)
reconstructed_output = decoder(encoded_repr)
autoencoder = keras.Model(input_images, reconstructed_output)

autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

autoencoder.fit(x_train, x_train)
```

```
# Encode and decode some digits
encoded_imgs = encoder.predict(x_test)
decoded_imgs = decoder.predict(encoded_imgs)
```




AE for image denoising (MNIST example)

Build the Model

```
input_layer = keras.Input(shape=(input_dim,))
encoding_layer1 = layers.Dense(128, activation='relu')(input_layer)
encoding_layer2 = layers.Dense(64, activation='relu')(encoding_layer1)
encoding_layer3 = layers.Dense(latent_dim, activation='relu')(encoding_layer2)
decoding_layer1 = layers.Dense(64, activation='relu')(encoding_layer2)
decoding_layer2 = layers.Dense(128, activation='relu')(decoding_layer1)
output_images = layers.Dense(input_dim, activation='sigmoid')(decoding_layer2)
autoencoder = keras.Model(input_layer, output_images)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

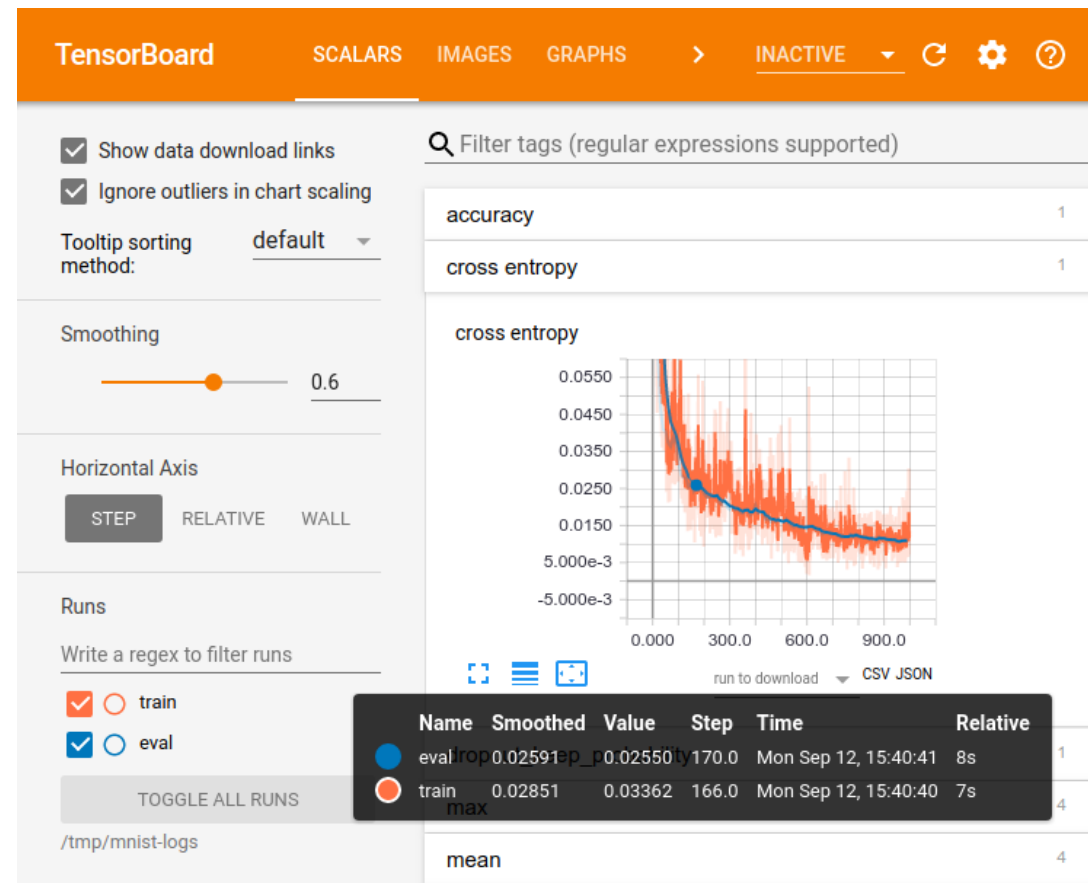


```
noise_factor = 0.4
x_train_noise = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noise = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

x_train_noise = np.clip(x_train_noise, 0., 1.)
x_test_noise = np.clip(x_test_noise, 0., 1.)
```



Keras: using callbacks

`keras.callbacks.ModelCheckpoint`
`keras.callbacks.EarlyStopping`
`keras.callbacks.LearningRateScheduler`
`keras.callbacks.CSVLogger`
`keras.callbacks.TensorBoard`



Start with high learning rate and then reduce it

- Predetermined piecewise constant learning rate
- Exponential scheduling

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.

Epoch 000,496 Learning rate 0.03 Activation Tanh Regularization None Regularization rate 0 Problem type Classification

DATA


Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10




REGENERATE

FEATURES

Which properties do you want to feed in?

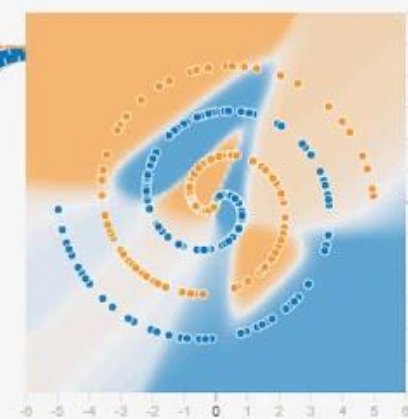


2 HIDDEN LAYERS



OUTPUT

Test loss 0.307
Training loss 0.208



Colors shows data, neuron and weight values.



☐ Show test data

☐ Discretize output